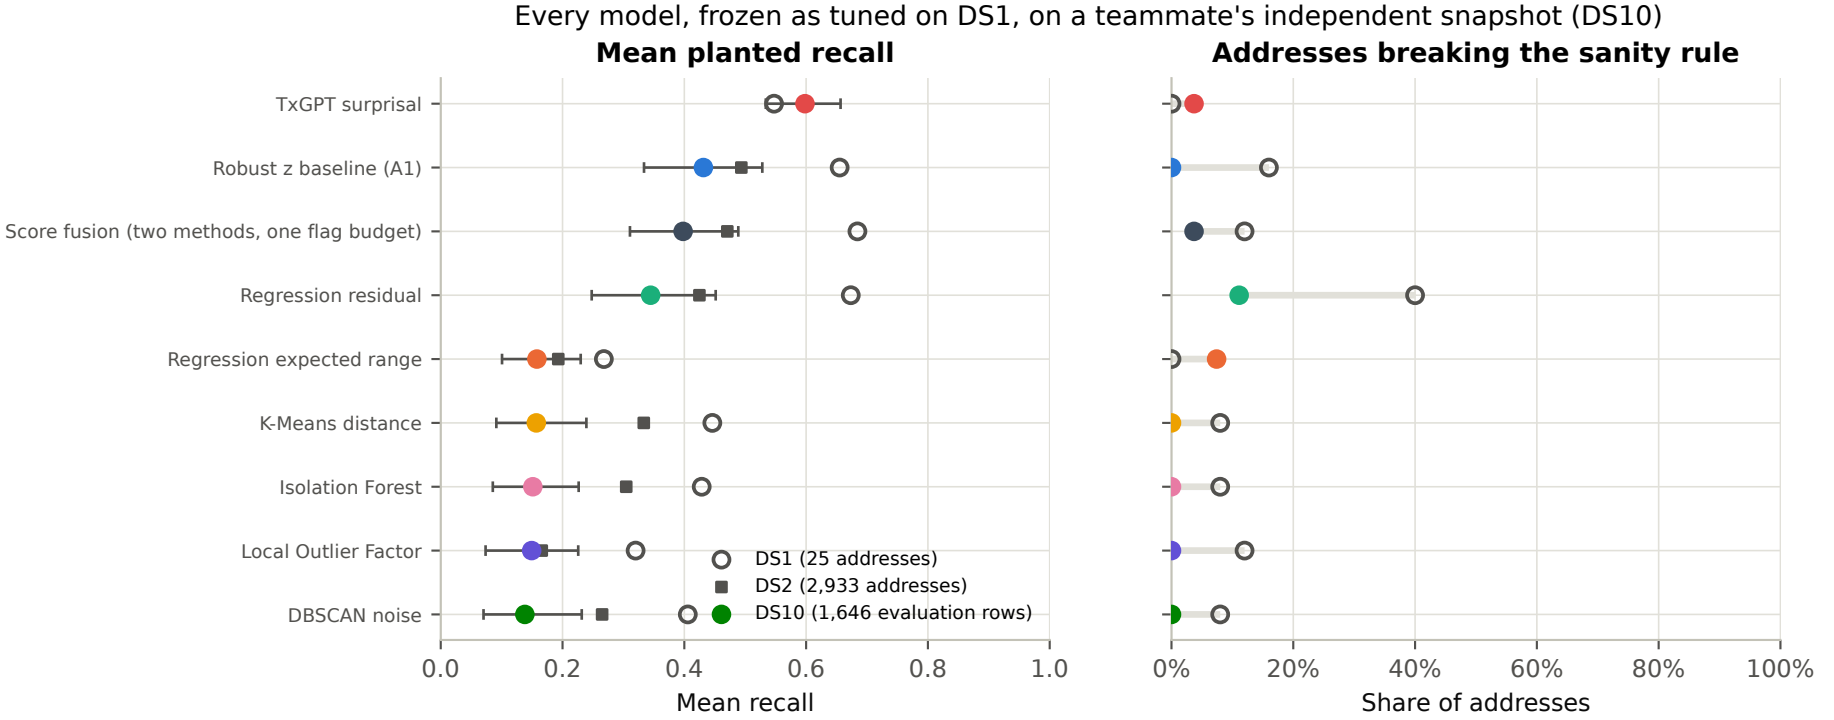
Every model, frozen as tuned on DS1, on a teammate's independent snapshot (DS10)
Mean planted recall
Addresses breaking the sanity rule
TxGPT surprisal
Robust z baseline (A1)
Score fusion (two methods, one flag budget)
Regression residual
Regression expected range
K-Means distance
Isolation Forest
Local Outlier Factor
DBSCAN noise
DS1 (25 addresses)
DS2 (2,933 addresses)
DS10 (1,646 evaluation rows)
0.0
0.2
0.4
0.6
0.8
1.0
Mean recall
0%
20%
40%
60%
80%
100%
Share of addresses