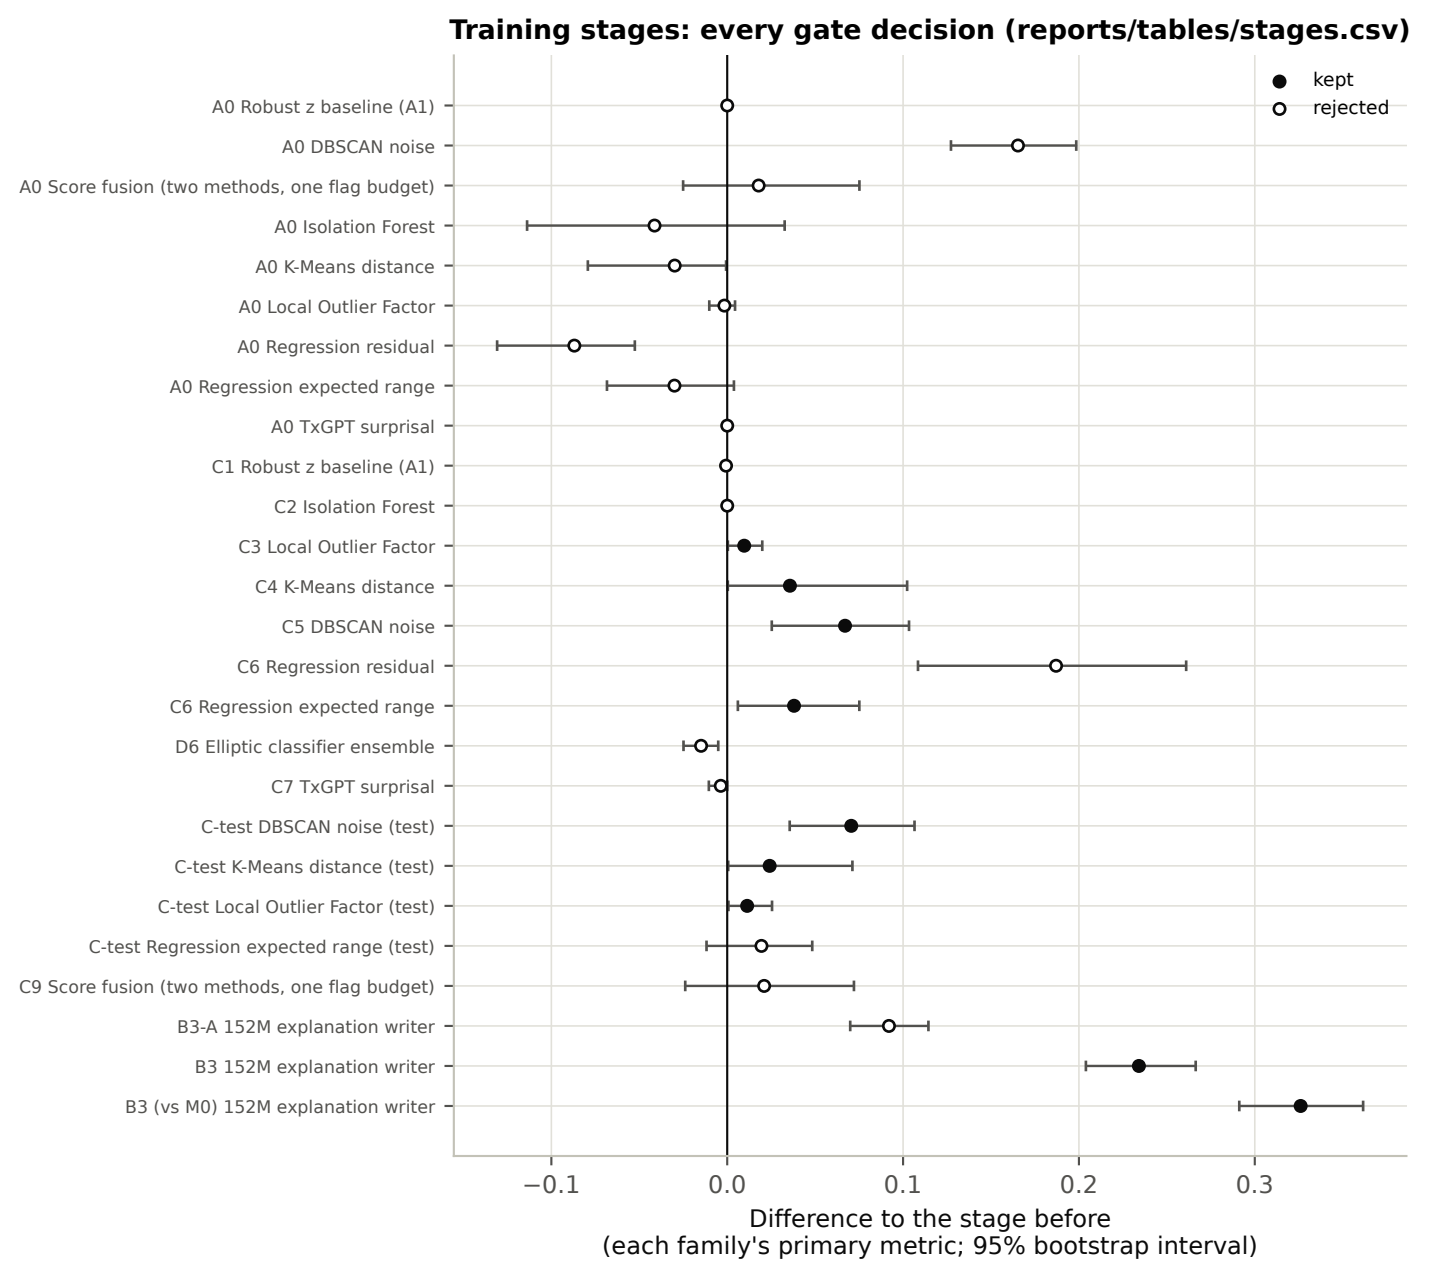
Training stages: every gate decision (reports/tables/stages.csv)
kept
rejected
A0 Robust z baseline (A1)
A0 DBSCAN noise
A0 Score fusion (two methods, one flag budget)
A0 Isolation Forest
A0 K-Means distance
A0 Local Outlier Factor
A0 Regression residual
A0 Regression expected range
A0 TxGPT surprisal
C1 Robust z baseline (A1)
C2 Isolation Forest
C3 Local Outlier Factor
C4 K-Means distance
C5 DBSCAN noise
C6 Regression residual
C6 Regression expected range
D6 Elliptic classifier ensemble
C7 TxGPT surprisal
C-test DBSCAN noise (test)
C-test K-Means distance (test)
C-test Local Outlier Factor (test)
C-test Regression expected range (test)
C9 Score fusion (two methods, one flag budget)
B3-A 152M explanation writer
B3 152M explanation writer
B3 (vs M0) 152M explanation writer
−0.1
0.0
0.1
0.2
0.3
Difference to the stage before
(each family's primary metric; 95% bootstrap interval)